


# Machine Learning for Streaming Data

IJCAI Tutorial 2024

Heitor Murilo Gomes<sup>1\*</sup>, Nuwan Gunasekara<sup>2</sup>  
Albert Bifet<sup>2,3</sup>, Bernhard Pfahringer<sup>2</sup>

\* Corresponding author: heitor.gomes@vuw.ac.nz

<https://nuwangunasekara.github.io/ijcai2024/>



[1] Victoria University of Wellington, New Zealand, [2] University of Waikato, New Zealand,  
[3] TELECOM Paris, LCTI, France.



# Heitor Murilo Gomes

Senior lecturer at the Victoria University of Wellington (VuW) in New Zealand. Before joining VuW, Heitor was co-director of the AI Institute at the University of Waikato. PI for a few research projects ranging from applied to fundamental research (i.e. ML for energy distribution, novel SSL approaches for DS, ...).

Leads the **capymoa** open source library for data stream learning, and provide support for **MOA** (Massive On-line Analysis).

<https://heitorgomes.com/>



# Nuwan Gunasekara

Completed his Ph.D. from University of Waikato in 2023. Currently works as a researcher at the AI Institute at the University of Waikato. Nuwan's research includes the development of new algorithms such as the Streaming Gradient Boosted Trees (SGBT), recurrent concept drifts and the intersection of Stream and Continual learning.

Nuwan is a core developer of **capymoa** and also provides support to **MOA**

[nuwangunasekara.github.io/](https://nuwangunasekara.github.io/)



# Albert Bifet

Professor of AI and the Director of the AI Institute at University of Waikato, and Professor of Big Data at Data, Intelligence and Graphs (DIG) LTCI, Télécom Paris, IP Paris. Co-chair of the NZ AI Researchers Association. Leads the TAIAO Environmental Data Science project and co-author of the book "Machine Learning from Data Streams" published at MIT Press.

Co-leads the open source project **MOA** Massive On-line Analysis, and provide advice/support for several other projects such as **capymoa**

<https://albertbifet.com/>



# Bernhard Pfahringer

Professor with the Department of Computer Science, and a co-director for the AI Institute, at the University of Waikato in New Zealand. Co-author of the book "Machine Learning from Data Streams" published at MIT Press. His research span a range of data mining and machine learning sub-fields, with a focus on streaming, randomization, and complex data.

Co-leads **MOA** and provide advice/support for **capymoa**

[www.cs.waikato.ac.nz/~bernhard/](https://www.cs.waikato.ac.nz/~bernhard/)



# About this tutorial

- **Our goal:** Introduce attendees to diverse machine-learning tasks for streaming data applications

# About this tutorial

- **Our goal:** Introduce attendees to diverse machine-learning tasks for streaming data applications
  - Classification, regression, ensemble learning, prediction intervals, concept drifts, partially and delayed streams, clustering, anomaly detection, ...

# About this tutorial

- **Our goal:** Introduce attendees to diverse machine-learning tasks for streaming data applications
  - Classification, regression, ensemble learning, prediction intervals, concept drifts, partially and delayed streams, clustering, anomaly detection, ...
- **Beyond the introduction:** Enable attendees to apply and extend the concepts introduced using python notebooks and *capymoa*



# Outline

- **Machine Learning for Streaming Data (intro)**  
[IJCAI\\_2024\\_introduction.ipynb](#)
  - Learning cycle
  - Evaluation procedures
  - Introduction to *capymoa*
- **Concept drifts**  
[IJCAI\\_2024\\_drifts.ipynb](#)
  - Simulation, Detection & Evaluation
- **Supervised Learning**  
[IJCAI\\_2024\\_supervised.ipynb](#)
  - Classification
  - Ensemble learning
- **Supervised Learning (cont.)**
  - Regression
  - Prediction Intervals  
[IJCAI\\_2024\\_prediction\\_intervals.ipynb](#)
- **Advanced Topics**  
[IJCAI\\_2024\\_advanced.ipynb](#)
  - Partially and delayed labeled streams
  - Clustering
  - Anomaly detection
  - More *capymoa* functionalities

**Notebooks:** [https://  
nuwangunasekara.github.io/  
ijcai2024/](https://nuwangunasekara.github.io/ijcai2024/)



# **Machine Learning for Streaming Data**

# Stream Learning

## What are data streams?

Sequences of items, possibly infinite, each item having a timestamp, and so a temporal order



# Stream Learning

## What are data streams?

Sequences of items, possibly infinite, each item having a timestamp, and so a temporal order

## Machine learning for streaming data (or Stream learning)

Data items arrive one by one, and we would like to **build and maintain models**, such as patterns or predictors, of these items in real time (or near real time)

# Stream Learning: Examples

Sensor data (IoT): energy demand prediction, environmental monitoring, traffic flow

Marketing and e-commerce: product recommendation, click stream analysis, sentiment analysis (social networks)

Cybersecurity: malware detection, spam detection, intrusion detection

And many more!\*

# Stream Learning

When should we abstract the data as a continuous stream?

# Stream Learning

When should we abstract the data as a continuous stream?

**can't store** all the data; or

**shouldn't store** all the data

# Stream Learning: **can't store**

Storing all the data may exceed the  
**available storage** capacity or cause  
practical limitations

The **volume or velocity** of incoming  
data may be too high to store and  
process in its entirety



# Stream Learning: **shouldn't store**

Storing all the data may not be desirable due to privacy concerns, compliance requirements, or the nature of the problem

For example, if we are only interested in **real-time analysis** or **immediate decision-making**

# Stream Learning

Using a stream abstraction, we can process the data incrementally, **focusing** on the **most recent** or **relevant** data points, and **discard** or **aggregate** the older data as needed


# Stream Learning

ML for **Batch** (“static”) data

**vs.**

ML for **Streaming** (“online”) data

# ML for Batch data



Fixed size dataset

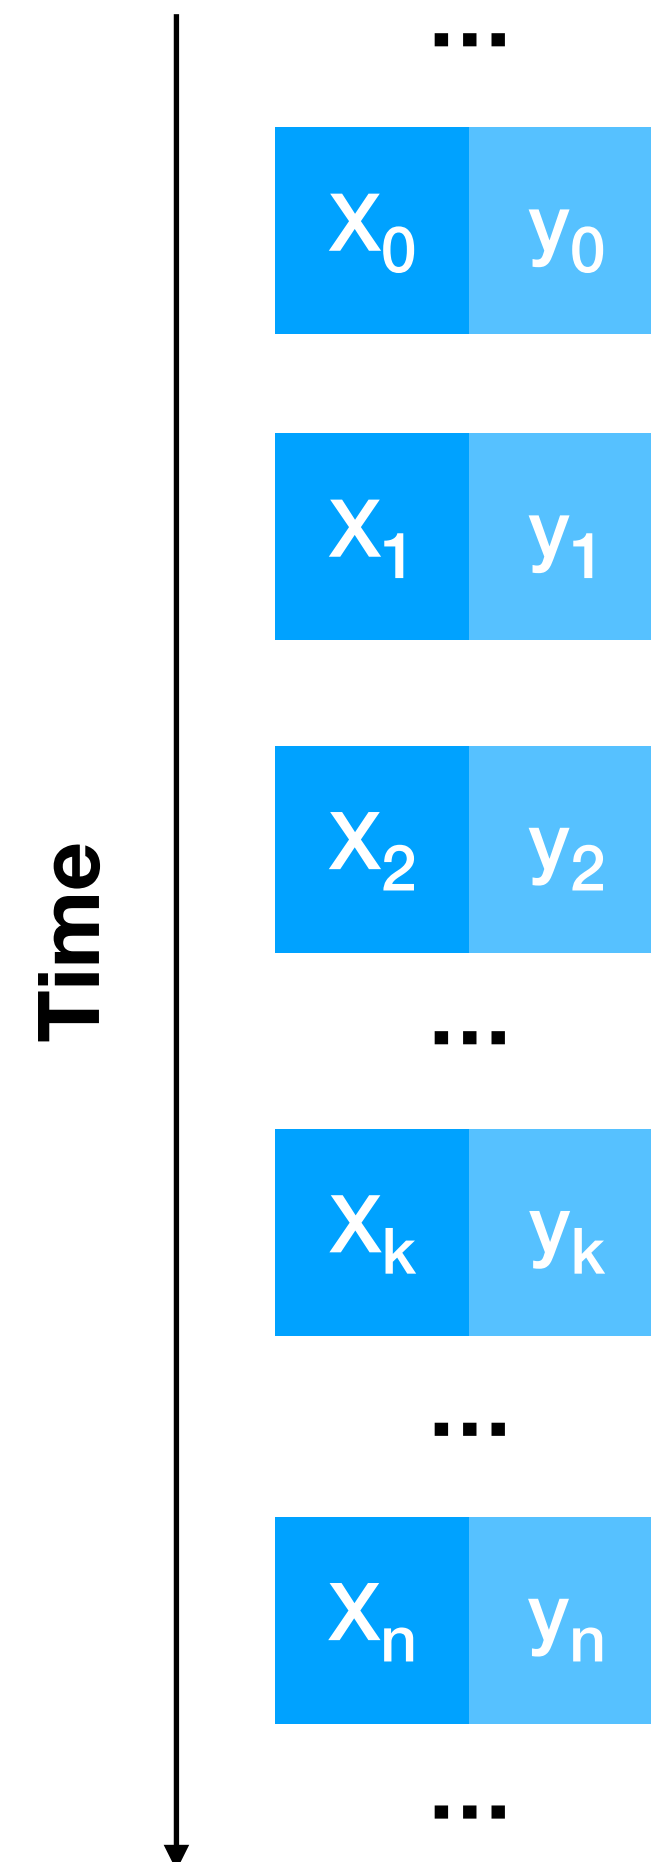
Random access to any instance

Well-defined phases  
(Train, Validation, Test)

## Challenges

noise, missing data,  
imbalance, high  
dimensionality, ...

# ML for Streaming data



Continuous flow of data

Limited time to inspect  
data points

Interleaved phases  
(Train, Validation, Test)

## Challenges

Concept drifts, concept  
evolution, strict memory/  
processing requirements,  
may more and...  
inherit all those from batch



# Batch vs. Streaming

## Batch data



The output is a **trained model**

## Streaming data



The output is a trainable model

# The Learning Cycle

# Batch vs. Stream

## Batch data

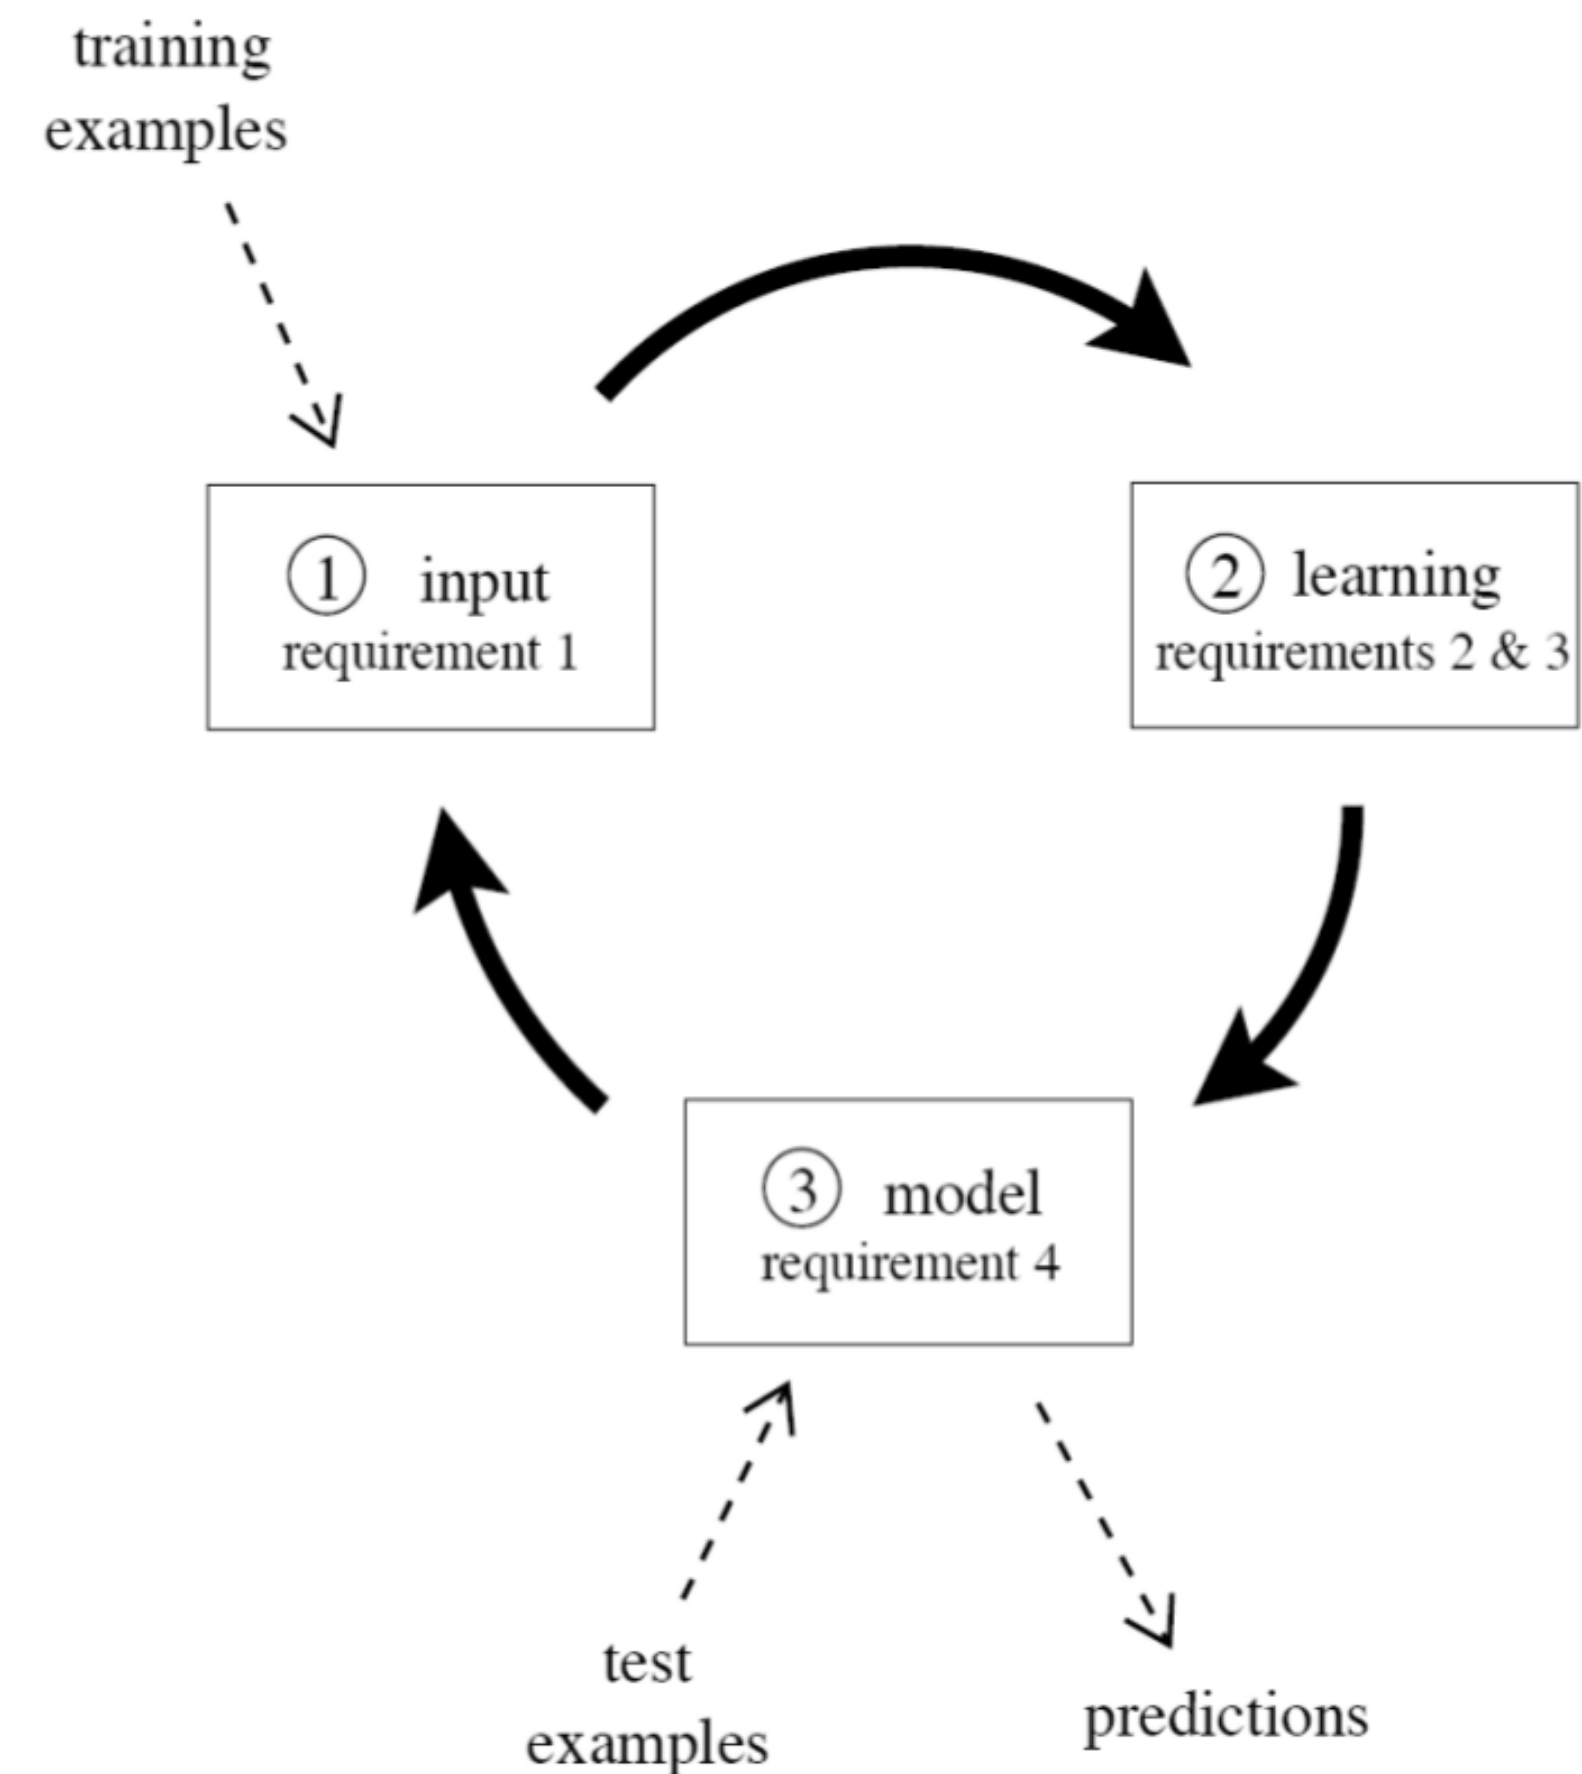
The model is updated through several passes over the training data

## Streaming data

The model is updated after observing every new data point

# The Learning Cycle

1. Process an example at a time, and **inspect it only once** (at most)
2. Use a **limited** amount of **memory**
3. Work in a **limited** amount of **time**
4. Be **ready to predict at any point**



# Evaluation



# Evaluation overview

Aspects concerning predictive performance evaluation:

- **Evaluation metrics.** How errors are considered?
- **Evaluation framework.** How past predictions influence the current metric?

Other measurements (e.g. wall-clock time, CPU time, ...)

# Evaluation Framework

## **(Periodic) Holdout**

- Interleave test and train subsets
- After testing in one window (or chunk), we observe the average over that window. We use the following window for training.

# Evaluation Framework

## Interleaved test-then-train (or cumulative)

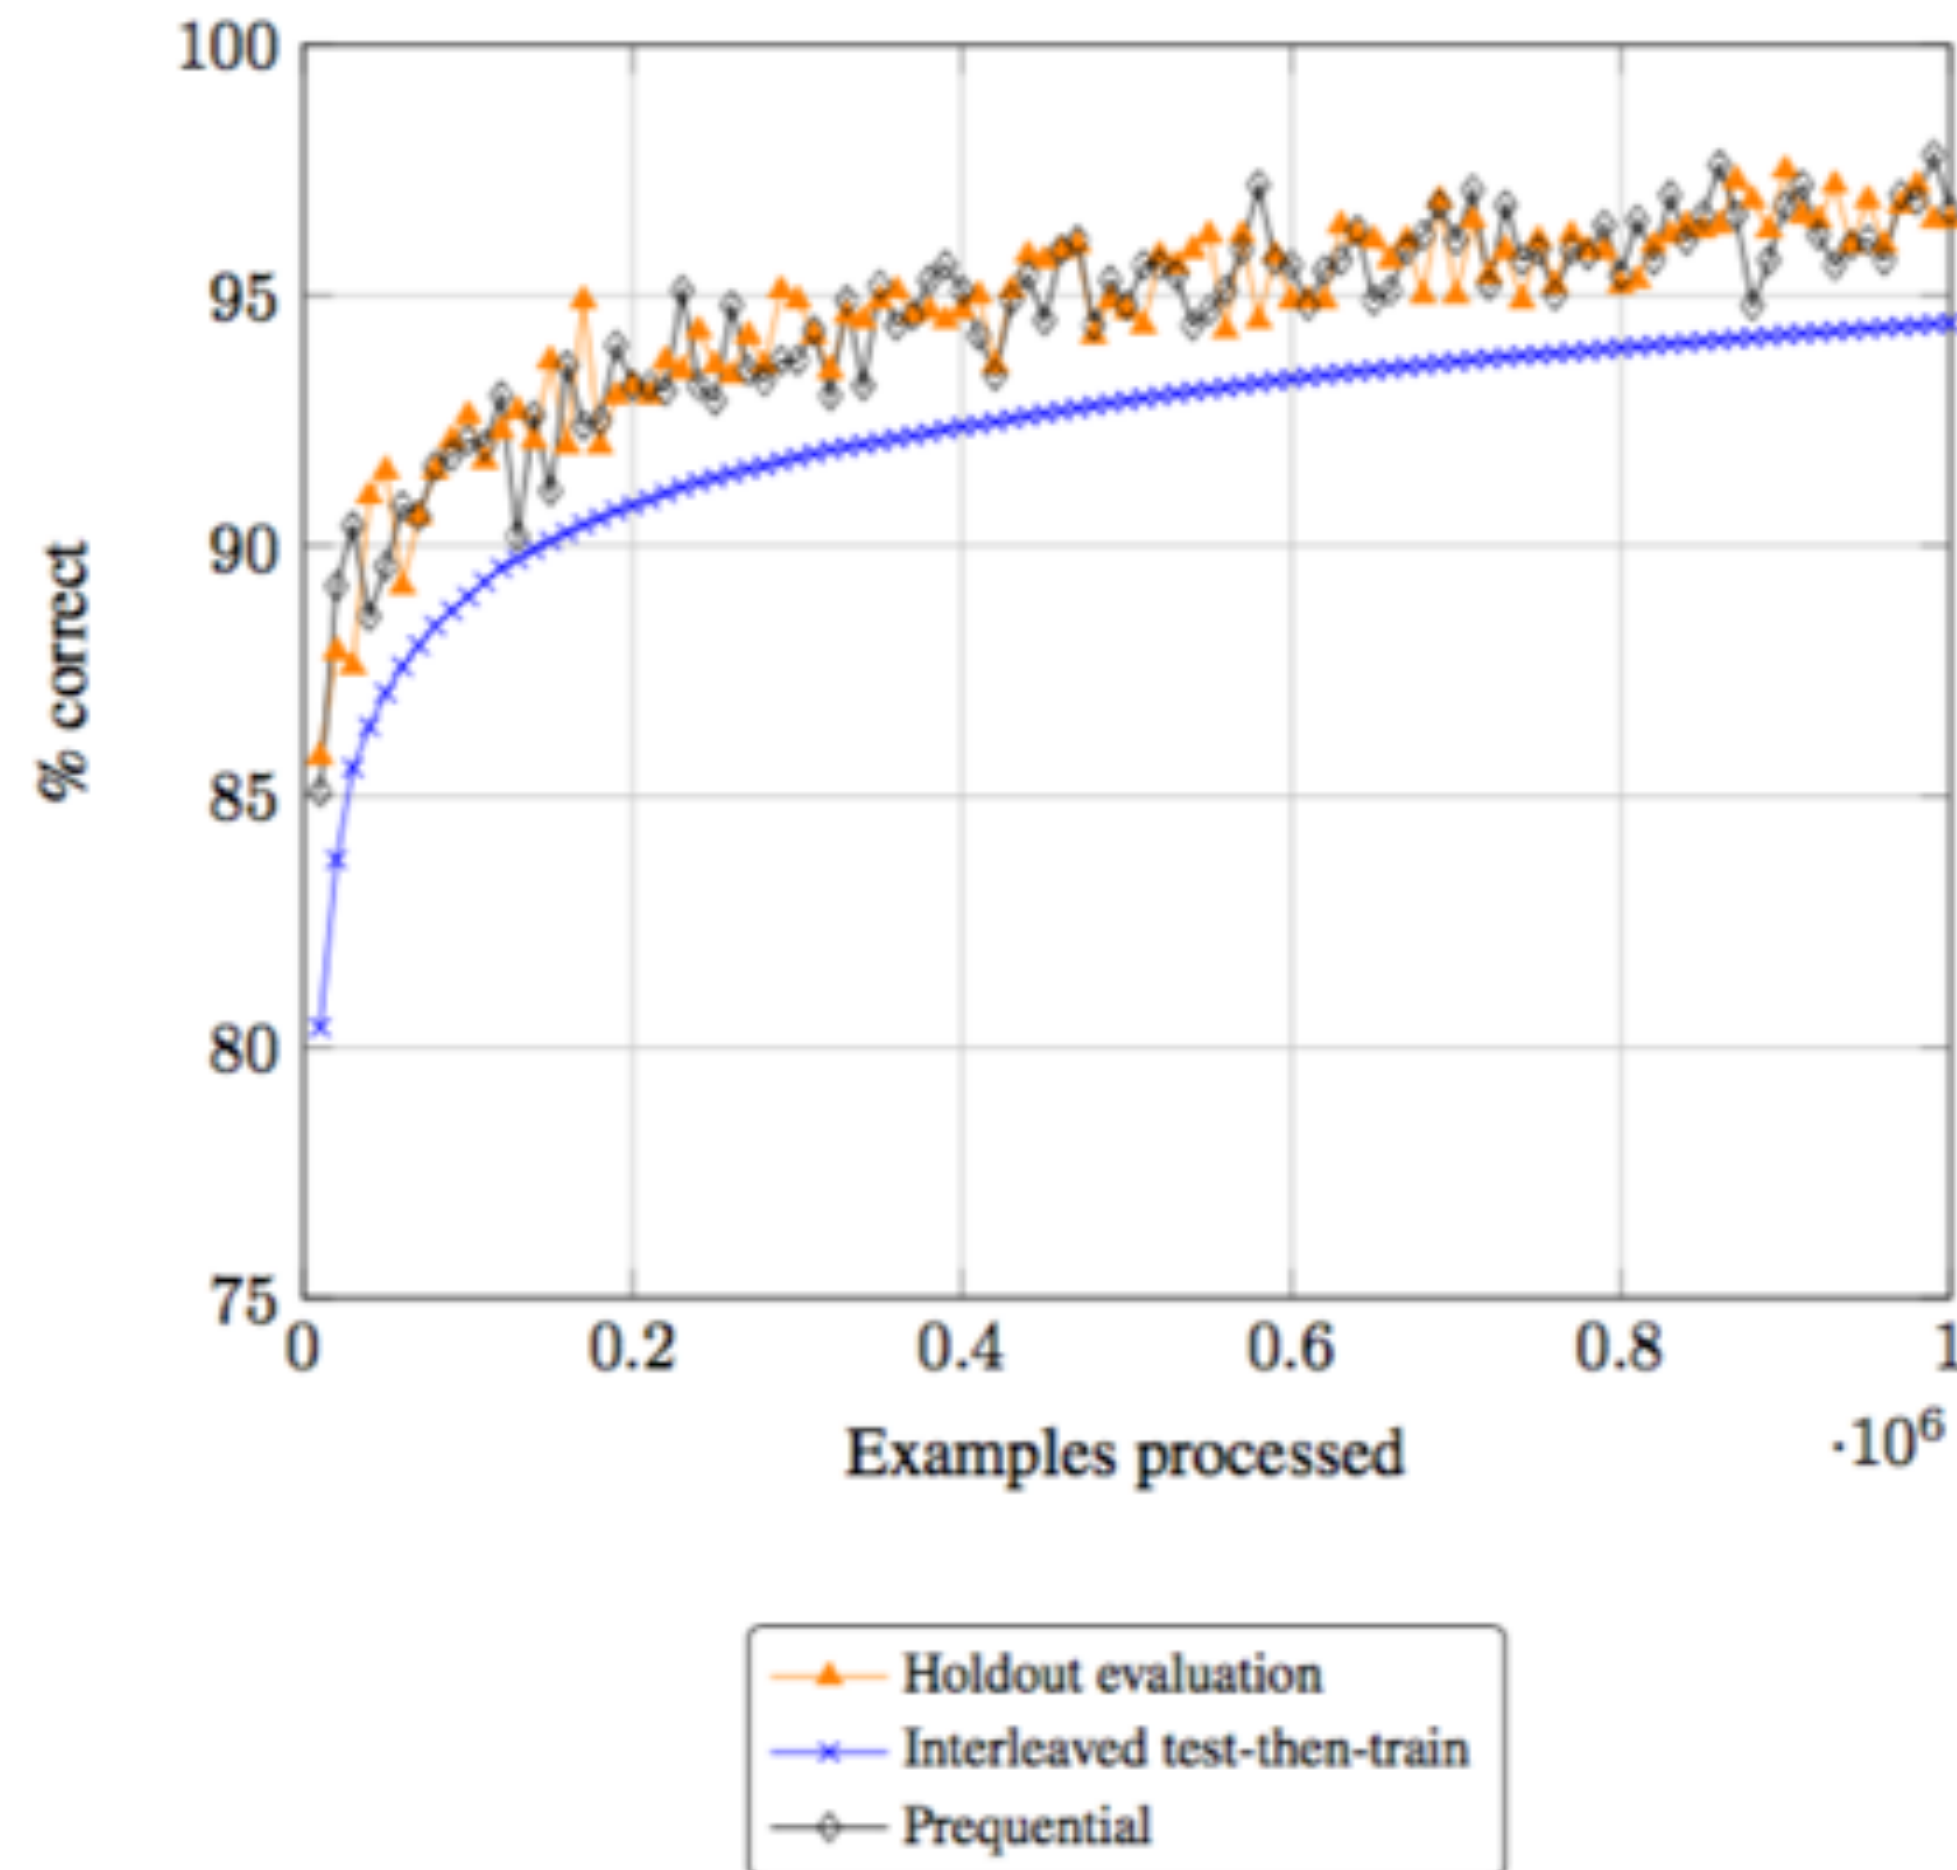
- Every instance is used for testing first, then for training
- At any point during execution, we observe the average over all instances seen so far

# Evaluation Framework

## Prequential evaluation

- Similar to interleaved test-then-train, but we observe the metrics over a sliding window of the latest instances (optionally, we can use a fading factor)

# Evaluation Framework






# CapyMOA

## Machine learning for data streams

<https://capymoa.org/>

<https://github.com/adaptive-machine-learning/CapyMOA>





# CapyMOA

A machine learning library for streaming data based on four pillars:

- **Efficiency**
- **Interoperability**
- **Accessibility**
- **Flexibility**

Other frameworks: **MOA** (java)<sup>1</sup>, **river** (python)<sup>2</sup> and **scikit-multiflow** (python)<sup>3</sup>

[1] Bifet, A., Holmes, G., Pfahringer, B., Kranen, P., Kremer, H., Jansen, T., & Seidl, T. (2010). Moa: Massive online analysis, a framework for stream classification and clustering. In *Workshop on applications of pattern analysis* (pp. 44-50). PMLR.

[2] Montiel, J., Halford, M., Mastelini, S.M., Bolmier, G., Sourty, R., Vaysse, R., Zouitine, A., Gomes, H.M., Read, J., Abdessalem, T. and Bifet, A., 2021. River: machine learning for streaming data in python. *Journal of Machine Learning Research*, 22(110), pp.1-8.

[3] Montiel, J., Read, J., Bifet, A., & Abdessalem, T. (2018). Scikit-multiflow: A multi-output streaming framework. *Journal of Machine Learning Research*, 19(72), 1-5.

# Why another one?

## Efficiency

A key aspect of stream learning are **efficient** implementations; they should learn from thousands of instances as quickly as possible (near real-time)

## Interoperability

Use algorithms from MOA, scikit-learn and PyTorch through an unified streaming API

## Accessibility


Must be easy to prototype experiments and extend functionality

## Flexibility


Advanced APIs for stream learning (concept drift, evaluation, visualisation, ...)

Code in either Python or Java\*

# Why? Efficiency



# Why? Efficiency



**Dataset:** RandomTreeGenerator with 2 abrupt drifts, 100k instances, 30 attributes, and 5 classes.

**Experiment:** 5 repetitions

**Algorithms:** NaiveBayes, HoeffdingTree (HT), Extremely Fast Decision Tree (EFDT), and AdaptiveRandomForest (the suffix indicates the number of base learners)

**Reproducibility:** <https://github.com/adaptive-machine-learning/CapyMOA/blob/main/notebooks/benchmarking.py>

# Why? Accessibility

Make it easier to configure and execute complex experiments

MOA's learning curve tends to be steeper as it is implemented in Java

CapyMOA allow users to code in Python, while allowing advanced users to take advantage from MOA objects directly from Python

```
from capymoa.stream.generator import SEA
from capymoa.stream.drift import DriftStream, AbruptDrift,
GradualDrift
from capymoa.classifier import AdaptiveRandomForestClassifier
from capymoa.evaluation import prequential_evaluation
from capymoa.evaluation.visualization import plot_windowed_results

SEA3drifts = DriftStream(stream=[SEA(1),
                                AbruptDrift(10000),
                                SEA(2),
                                GradualDrift(start=20000,
                                              end=25000),
                                SEA(3),
                                AbruptDrift(45000),
                                SEA(1)])

arf =
AdaptiveRandomForestClassifier(schema=SEA3drifts.get_schema(),
                               ensemble_size=100,
                               number_of_jobs=4)

results = prequential_evaluation(stream=SEA3drifts,
                                learner=arf,
                                window_size=1000,
                                max_instances=50000)

print(f"Cumulative accuracy = {results['cumulative'].accuracy()}")
print(f"wallclock = {results['wallclock']} seconds")
display(results['windowed'].metrics_per_window())
plot_windowed_results(results, ylabel='Accuracy')
```



# Why? Interoperability

Easy access to **MOA**, **PyTorch** and **Scikit-learn** learners

Wraps and expose MOA's API through a modern and simple API

Standardise evaluation and benchmarking without losing **flexibility**



# Why? Flexibility

CapyMOA facilitates preparing and executing experiments by incorporating:

- High-level evaluation functions
- Build-in visualisation tools
- Advanced Concept Drift API; and more
- We will show several examples throughout this tutorial



# CapyMOA team

- Heitor Murilo Gomes (project leader)<sup>1</sup>
- Anton Lee<sup>1</sup>
- Nuwan Gunasekara<sup>2</sup>
- Yibin Sun<sup>2</sup>
- Guilherme Cassales<sup>2</sup>
- Marco Heyden<sup>3</sup>
- Justin Liu<sup>2</sup>
- Jesse Read<sup>4</sup>
- Maroua Bahri<sup>5</sup>
- Marcus Botacin<sup>6</sup>
- Vitor Cerqueira<sup>7</sup>
- Albert Bifet<sup>2,9</sup>
- Bernhard Pfahringer<sup>2</sup>
- Yun Sing Koh<sup>8</sup>

And many other individual contributors

[1] Victoria University of Wellington, New Zealand

[2] University of Waikato, New Zealand

[3] KIT, Germany

[4] École polytechnique, IP Paris, France

[5] INRIA Paris, France

[6] Texas A&M Engineering, USA

[7] Porto University, Portugal

[8] University of Auckland, New Zealand

[9] Télécom Paris, IP Paris, France





# CapyMOA summary

- Allows access to existing and future MOA implementations
- Code in Python or Java, or combine both (e.g. Python using MOA objects)
- Minimal overhead in comparison to executing native MOA
- Integration with PyTorch and scikit-learn
- Streams, learners and evaluation are designed to interoperate with visualisation
- 3 releases (Mar/24, May/24, July/24)
- 20 classifiers, 8 regressors, 11 drift detectors, 3 anomaly detectors, ... as of 0.6.0



# Practical examples

IJCAI\_2024\_introduction.ipynb